

What does one country review cost?

National uptake of Recommendation Rec(2022)2
Budgeting note for Micheal's / CEDS team
28 September 2026

Plain answer

For a Belgium-like desk review, plan on roughly 5 million billed tokens (mostly input: the agent re-reads a growing context as it searches and writes).

At current Cursor list prices (API rates), **without counting your included allowance**, that is roughly:

Model band	Approx. cash cost per country	In pounds (at £0.75 / \$1)
Cheap / fast (Composer 2.5, Gemini Flash, GPT Luna)	about \$1–\$4	about £0.80–£3
Mid (Grok 4.6/4.7, Claude Sonnet 5, GPT Terra)	about \$10–\$11	about £7.50–£8.50
Strong frontier (Claude Opus 5 / 5.5, GPT-5.6 Sol)	about \$21–\$26	about £16–£20
Top / expensive (Claude Fable 5.1)	about \$50+	about £40+

Inside Cursor: that usage usually draws your **included monthly allowance first**. You only see a separate invoice line if you go on-demand after the pool is used up. Cloud Agents are billed at the **selected model's API price**; there is no separate charge for the VM itself.

These numbers are **planning estimates**, not the Cursor invoice. Exact spend is on the Cursor usage dashboard for the run.

How we estimated tokens (Belgium PoC)

This run produced the Belgium desk review (countries/belgium.md + .yaml) as the first task.

We reconstructed billed tokens from the saved agent transcript (224 messages: 6 user, 141 assistant, 77 tool), using a cumulative-context model:

- each model call is charged for the **context so far** (input) plus the **new assistant message** (output, including tool-call arguments)
- ~8k tokens of system/tool-schema overhead assumed per call

Slice	Model calls (approx.)	Input tokens	Output tokens	Total
Belgium desk review only (to first follow-up)	~80	~4.9 million	~0.01–0.05 million	~5 million
Whole session so far (Belgium + method PDFs + Q&A)	~140	~13.6 million	higher	~14 million

Planning figure used below: 5.0M input + 0.05M output for one country desk review (Belgium-like).

Why input dominates: agents re-send prior tool results (web pages, laws, drafts) on later turns. The written brief is “only” ~50 KB of files; the **search trail** is what costs tokens.

Planning bands (tokens per country)

Band	When it happens	Tokens to budget
Lean	English-heavy sources, few instruments, little back-and-forth	~1.5M input + 0.03M output
Belgium-like (baseline)	Dedicated statute + decrees, FR/NL/EN mix, full 58 indicators	~5M input + 0.05M output

Band	When it happens	Tokens to budget
Hard	Multi-jurisdiction (e.g. Länder), many association rulebooks, several human follow-ups	~12M input + 0.12M output

Cost per model — Belgium-like country (~5M in / 0.05M out)

Prices from [Cursor Models & Pricing](#) (per million tokens, USD).

GBP converted at **£0.75 per \$1** (planning FX only; check the day's rate).

Assumes **no prompt-cache discount** and **no Teams “Cursor Token Rate”** (see caveats).

Model	Provider pool	Input \$/MTok	Output \$/MTok	Est. USD / country	Est. GBP / country	Pros	Cons
Composer 2.5	Cursor Models	0.50	2.50	~\$2.60	~£2.00	Cheapest Cursor-native option; good for drafting; draws Cursor Models pool	Weaker on hard legal nuance / long multilingual research than frontier models
Gemini 3.8 Flash	Other Models	0.75	3.50	~\$3.90	~£2.90	Fast and cheap; fine for first pass	May miss subtle statutory distinctions; more human checking needed
GPT-5.6 Luna	Other Models	0.20	1.20	~\$1.10	~£0.80	Lowest cash cost in this table	Smallest / least capable GPT-5.6 tier — risky as sole reviewer for legal desk work
Grok 4.6 / 4.7	Cursor Models	2.00	6.00	~\$10.30	~£7.70	Strong agent use; Cursor Models pool; no Teams token surcharge on first-party	Still not a specialist legal model; long-context / Fast modes cost more
Claude Sonnet 5	Other Models	2.00	10.00	~\$10.50	~£7.90	Best cost/quality balance for this task type; solid long docs	Other Models pool; Teams add Cursor Token Rate
GPT-5.6 Terra	Other Models	2.00	12.00	~\$10.60	~£8.00	Mid GPT-5.6; capable agent	More expensive output than Sonnet; Other Models pool
Claude Opus 5.5	Other Models	4.00	20.00	~\$21.00	~£15.80	Stronger reasoning / careful reading	~2× Sonnet cost
GPT-5.6 Sol	Other Models	4.00	20.00	~\$21.00	~£15.80	Strong OpenAI frontier (promo pricing noted by OpenAI into Nov 2026)	Same ballpark as Opus 5.5; burns Other Models pool faster
Claude Opus 5	Other Models	5.00	25.00	~\$26.30	~£19.70	Very strong for dense legal mapping	Pricey for 30+ countries
Claude Fable 5.1	Other Models	10.00	50.00	~\$52.50	~£39.40	Top-tier long-running agent intelligence	~5× Sonnet; needs retention approval in some setups; usually overkill for routine scans

Lean and hard countries (same models, quick view)

Model	Lean (~1.5M/0.03M)	Belgium-like (~5M/0.05M)	Hard (~12M/0.12M)
Composer 2.5	~£0.60	~£2.00	~£4.70
Gemini 3.8 Flash	~£0.90	~£2.90	~£7.10
Grok 4.6/4.7	~£2.40	~£7.70	~£18.50
Claude Sonnet 5	~£2.50	~£7.90	~£18.90

Model	Lean (~1.5M/0.03M)	Belgium-like (~5M/0.05M)	Hard (~12M/0.12M)
Claude Opus 5.5 / GPT-5.6 Sol	~£5.00	~£15.80	~£37.80
Claude Opus 5	~£6.20	~£19.70	~£47.30
Claude Fable 5.1	~£12.40	~£39.40	~£94.50

Scale-up sketch (Belgium-like, Claude Sonnet 5)

Volume	Est. tokens	Est. cash if no allowance	Notes
1 country	~5M	~£8	PoC
10 countries	~50M	~£80	Still small vs human days
30 States Parties	~150M	~£240	Plus human expert time (the real cost)
30 hard countries	~360M	~£570	Upper planning bound

Human expert review time is **not** in these figures and will usually dwarf token cost.

What Cursor actually charges (so the allowance point is clear)

1. **Cloud Agents bill tokens at the model's API list price** (same rates as in the table).
2. **Included allowance is used first** (Cursor Models pool vs Other Models pool, depending on model).
3. **On-demand** applies only after included usage is exhausted (must be enabled for Cloud Agents).
4. **No separate VM fee** for Cloud Agents (as of Cursor's public docs / forum clarifications).
5. On **Teams / Enterprise**, third-party models also incur a **Cursor Token Rate of \$0.25 per million tokens** on top of API price (~+£0.90 on a 5M-token Sonnet run). First-party Cursor models (Grok, Composer) are exempt.
6. **Prompt caching** can cut effective input cost a lot when the same law texts are re-read; the table above is the **pessimistic (no-cache)** case.

So: **"How much does a country cost?"**

- **~5M tokens** for Belgium-like work.
- **~£2 to ~£20** cash-equivalent depending on model, if it were all on-demand.
- **Often £0 extra on the invoice** while it still fits inside included allowance — but it still **consumes** that allowance.

Practical recommendation for this project

Approach	Suggestion
Default for production desk reviews	Claude Sonnet 5 or Grok 4.6/4.7 — about £8 / country cash-equivalent, good quality for statute mapping
Budget / volume pass	Composer 2.5 or Gemini Flash first, then human or Sonnet spot-check of partial / not_found lines
Hard countries or expert-facing final polish	Opus 5.5 / GPT-5.6 Sol for a second pass on weak indicators only (not the whole 58 from scratch)
Do not default to Fable	Reserve for cases where Sonnet/Grok clearly fail long-horizon research

Always keep the **human check** step (see `how-a-country-scan-works.pdf`): token cost is cheap compared with sending a wrong `met` to a country expert.

Caveats (read before budgeting)

- Token estimate is **reconstructed from this PoC transcript**, not copied from the Cursor billing API. Check **Settings** → **Usage** for the true number on your account.
 - Output tokens in the transcript look low because much “writing” sits in tool-call payloads; we padded output to **0.05M** for Belgium-like planning.
 - Follow-up chat (method notes, PDFs, Q&A) can **double or triple** a session beyond the pure country file. Budget **desk review** and **partner communication** separately if needed.
 - FX (£0.75/\$1) is illustrative.
 - Prices change; re-check cursor.com/docs/models-and-pricing when locking a grant budget.
 - Web search / browsing tool usage may add provider-side costs on some stacks; Cloud Agent public guidance is still “tokens at API rates.”
-

One-line budget line for a proposal

AI desk review: ~5 million tokens per country (~£8 at Claude Sonnet 5 list rates); usually drawn from Cursor included allowance; human expert validation separate.

Prepared from the Belgium PoC run (bc-066382ea-3141-4635-b3da-bec7d66c6022) and Cursor published model prices as of 28 September 2026.